



# **Intelligent IIoT Service Deployment using MEC and Federated Learning in Drone- Assisted Smart Agriculture**

**Team No. 16  
EDGLERS**



# PROBLEM STATEMENT

- **Modern farming faces challenges in rapidly deploying and managing digital services like crop monitoring, pest detection, and irrigation control.**
- **Traditional cloud-based or manual IT setups cause high delay (latency), network congestion, and integration obstacles—especially over vast or remote fields where reliable connectivity is critical.**
- **Manual or slow deployment makes it impossible to react to sudden changes (like weather or pests) and to scale up for large harvest operations.**





# Topology & Modules

## Topology Design

- Cloud (Node 0): powerful central server (analytics).
- MEC/Fog (Node 1): local farm server, mid-compute, low latency.
- Drones (Nodes 50+): mobile collectors, Wi-Fi & mesh links.
- Soil Sensors (Nodes 100+): low-power IoT devices, LoRaWAN links + backup to MEC.

## Applications/Modules

- SensorDataCollection: sensors create soil readings.
- Imagecapture: takes photos from csv file and feeds it to the model
- DroneProcessing: drones process & forward data to MEC.
- CloudAnalytics: MEC forwards aggregated data to cloud.
- Placement:
  - Sensors → sensor nodes
  - Drones → drone nodes
  - MEC → node 1
  - Cloud → node 0



```
def scheduleTask(self, sensorId, currTime):
    sensorPos = self.sensorPositions[sensorId]
    bestDrone, minCost = None, float('inf')

    for droneId, dronePos in self.dronePositions.items():
        dist = self.calculateDistance(dronePos, sensorPos)
        batPenalty = (100 - self.droneBattery[droneId]) * 0.5
        cost = dist + batPenalty

        if self.droneBattery[droneId] > 20 and cost < minCost:
            minCost, bestDrone = cost, droneId

    if bestDrone is not None:
        self.assignedTasks[bestDrone].append({...})
        self.dronePositions[bestDrone] = sensorPos
        # Decrease battery (min 5 units, or remaining battery if <5)
        self.droneBattery[bestDrone] -= min(5, self.droneBattery[bestDrone])
        return bestDrone

    return None
```

# Task Scheduling & Simulation Logic

## Task Scheduling (Greedy Algorithm)

- $\text{Cost} = \text{distance}(\text{drone}, \text{sensor}) + 0.5 \times (100 - \text{battery}\%)$
- Only drones with  $>20\%$  battery considered.
- Best drone chosen  $\rightarrow$  moves to sensor, collects data, battery  $-5\%$ .
- If battery  $<30\% \rightarrow$  upload data to MEC & recharge  $(+20\%)$ .

## Simulation Logic

- Event generation:
  - Sensors  $\rightarrow$  every 300 units, Drones  $\rightarrow$  every 250 units
  - MEC  $\rightarrow$  sends to cloud every 800 units
- Custom Strategy: runs periodically to assign drones dynamically.
- Outputs (CSV):
  - Soil sensor data, Drone collection logs, MEC processing results, YAFS trace for latency & paths

# Questions from Review one

**1.How does the proposed solution enable fast reaction to sudden environmental changes?**

Drones perform local training on new data, enabling rapid self-evolution and immediate adaptation to localized environmental changes without waiting for central commands.

**2.How do you scale when multiple drones/sensors are working together?**

The system scales efficiently as each Drone Client works independently, with the server only collecting model weights. Asynchronous Federated Averaging avoids bottlenecks, maintaining performance as the fleet grows.

**3.What happens if one drone or MEC node fails—does the whole system stop?**


The system is resilient because its decentralized design lets other drones keep working even if one fails. The aggregation script uses available updates, so the global model keeps improving.

## Changes from review 1

- 1.Changes the LoRaWAN receiver from independent drones to the MEC server
- 2.Added an FL framework that updates the drone models every month

## Tasks left for review 3

- 1.Performance analysis yet to be implemented on the cloud server
- 2.Complete model and architecutre integration
- 3.Yeild prediction algorithms and models should be implemented





# YAFS Simulation Framework Implementation

01

## Simulation Components

- **Topology Builder:** Creates hierarchical network with nodes (sensors, drones, MEC, cloud) and edges (LoRaWAN, WiFi, fiber connections)
- **Application Modeling:** Defines data flow from sensors → drones → MEC → cloud with realistic message sizes and processing requirements
- **Resource Management:** Simulates CPU, RAM, storage, and network bandwidth usage across all nodes

02

## YAFS Features Utilized

- **Dynamic Message Routing:** Realistic network latency and bandwidth simulation\
- **Resource Monitoring:** Real-time tracking of node utilization and performance metrics
- **Event-Driven Simulation:** Time-stepped execution modeling actual network behavior

# CNN Model For Pest Detection

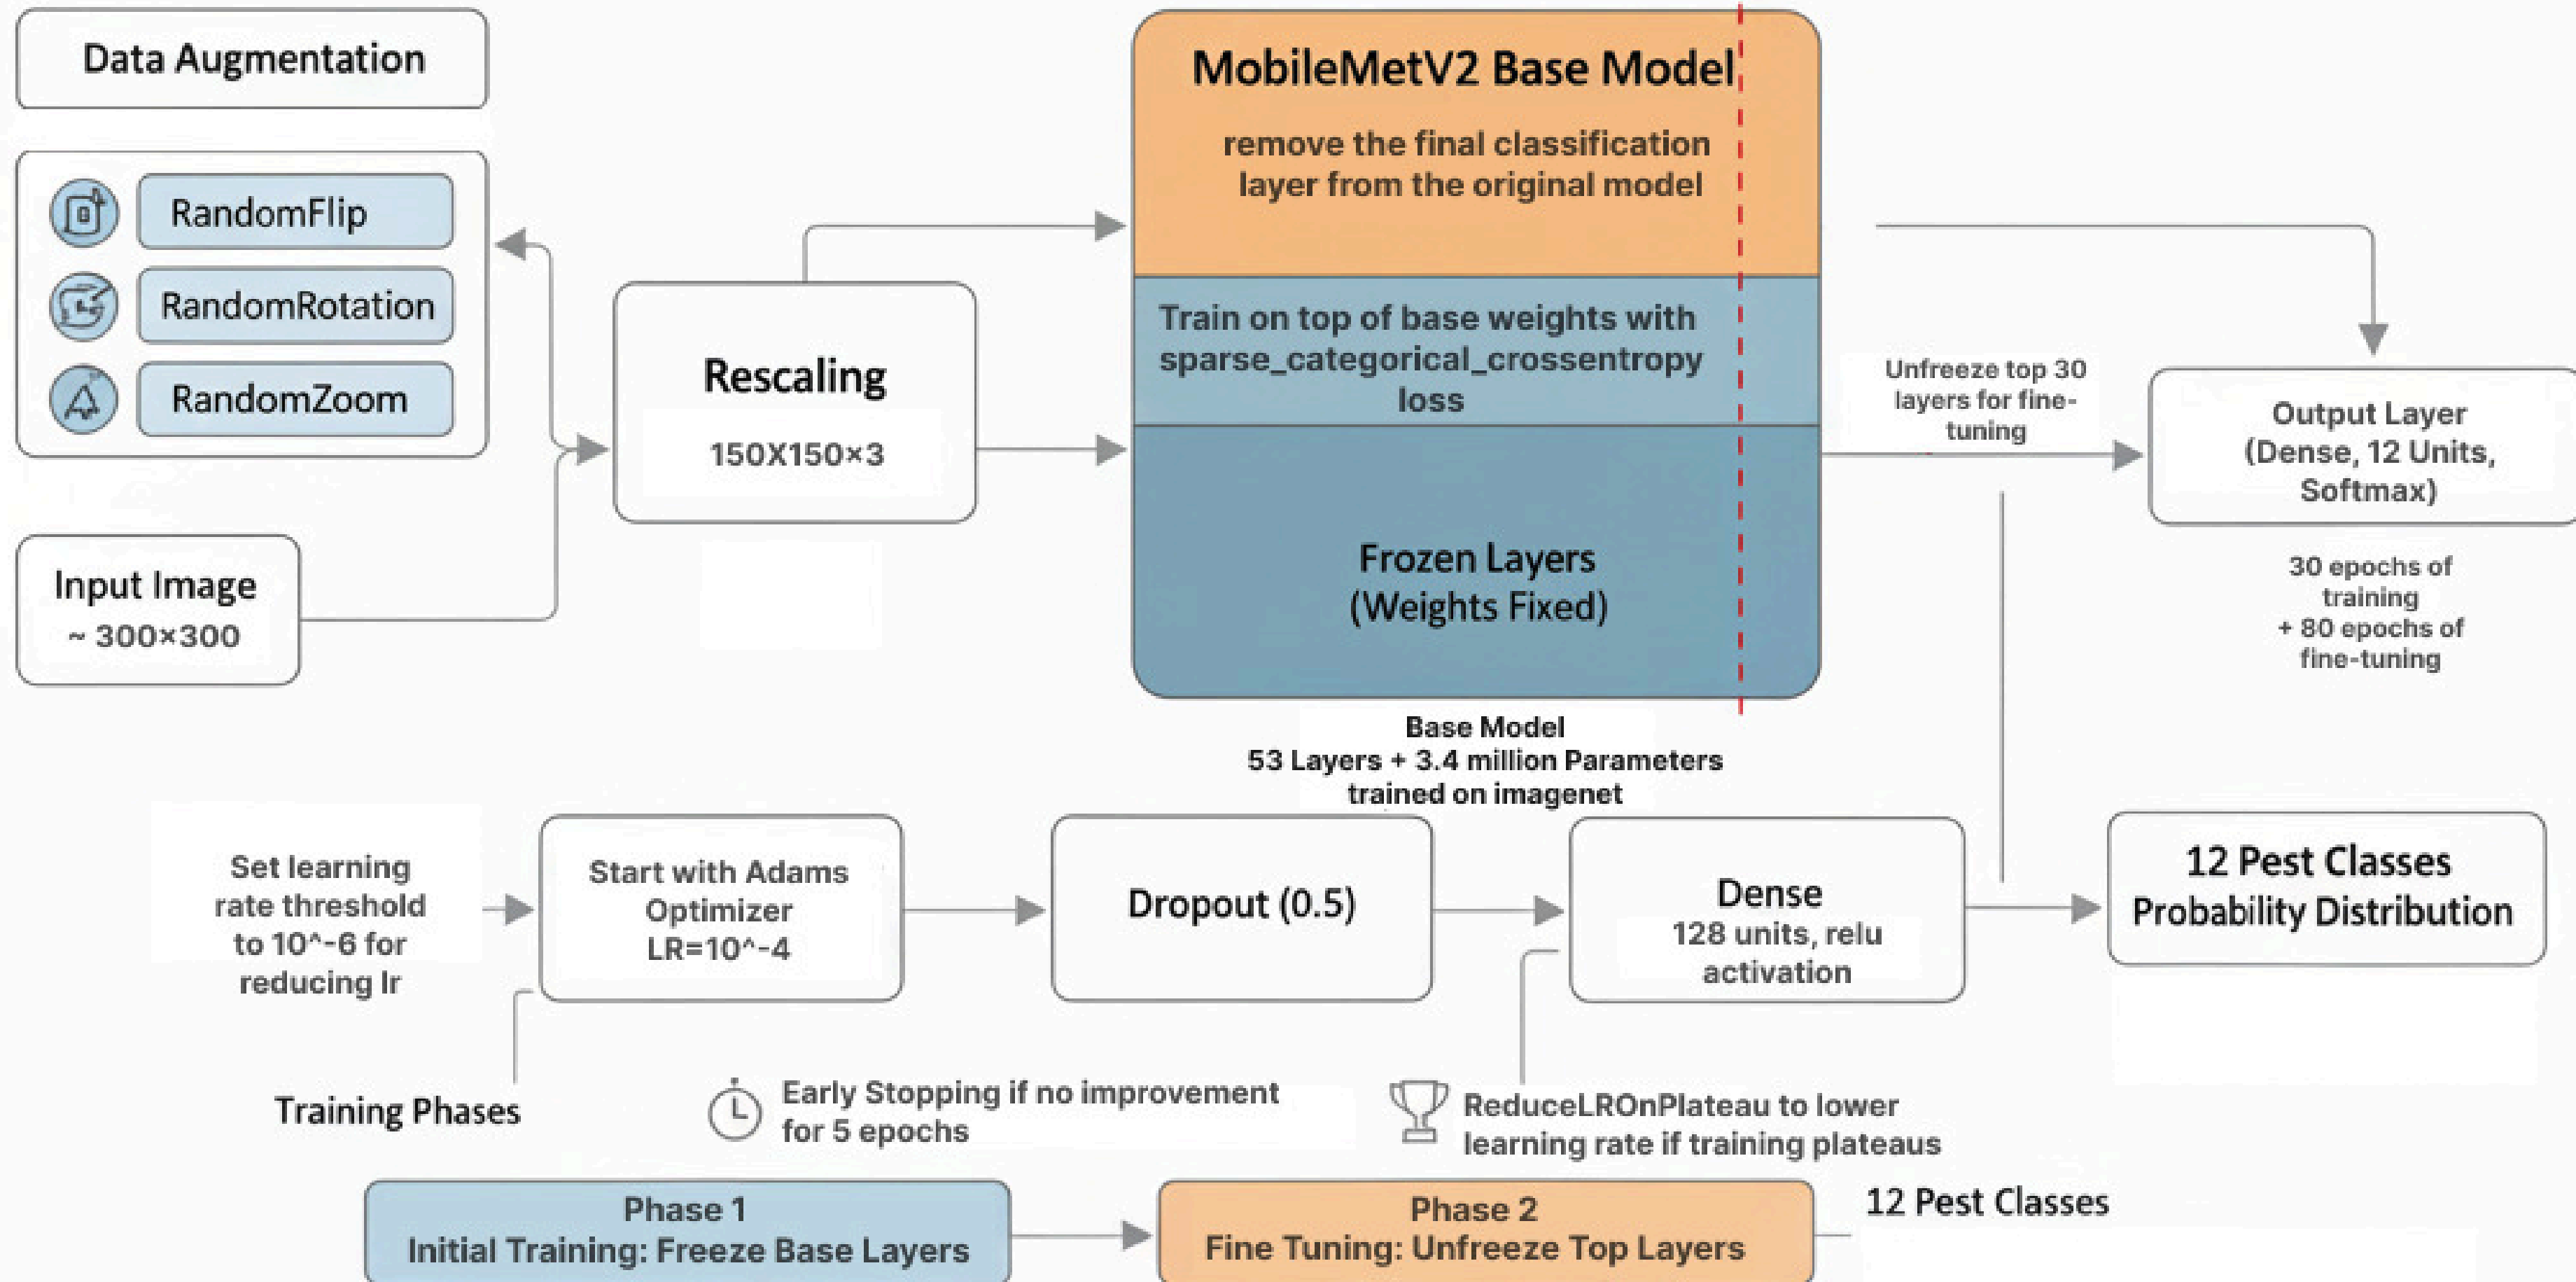
## Initial Attempt: Built a CNN from scratch:

- Resized images ( $150 \times 150 \times 3$ ), 80-20 split, data augmentation
- Architecture: Conv2D + MaxPooling → Dense layers → Softmax (12 pest classes)
- Trained with Adam optimizer, sparse categorical cross-entropy
- Achieved ~50% accuracy with ~60MB model (20-30 epochs, early stopping)

## Second Iteration: Applied Transfer Learning with MobileNetV2:

- Used pretrained ImageNet weights, removed top layers
- Added Global Average Pooling, Dropout(0.3), Dense(12, Softmax)
- Used callbacks: Early Stopping + Reduce LR on Plateau
- Fine-tuned by unfreezing last 30 layers, retrained with lower LR
- Achieved ~85% accuracy with ~110MB model (110 epochs, early stopping)

# System Architecture CNN Model



# Federated Learning Model

Built a Federated Learning system that trains a new neural network taking weights from CNN models from drone fleet, improving the model's pest detection accuracy.

## How It Works

- Initially a central server distributes a global Neural Network to drone clients.
- Each drone improves the model locally on its private data.
- Drones upload only their small model weight updates, which the server combines using Federated Averaging (FedAvg) to create an improved global model.

Designed a 3-part FL simulation that trains a model across distributed devices, solving the key challenge of network efficiency.

## Result & Verification:

The final model achieved ~98% accuracy on a test dataset, proving the FL system is a high-performance, privacy-preserving solution.

# Light-weight Fuzzy Logic for Soil pH Classification

## Objective

- To build a lightweight, interpretable CI model that can run on a low-power microcontroller for real-time, on-site soil analysis.

## How It Works

- The model uses a Fuzzy Inference System with a simple "if-then" rule base.
- It takes a precise pH value from a sensor and classifies it into one of three linguistic categories: Acidic, Neutral, or Alkaline.
- This approach demonstrates a rule-based CI technique, which is highly efficient and easy to understand.

The final working script achieved ~97% accuracy in classifying soil pH.

## Future Work: Build a Adaptive Neuro-Fuzzy Inference System (ANFIS)

To build a hybrid CI model that combines the learning capabilities of a Neural Network with the transparent, rule-based structure of a Fuzzy Logic system



```

[starkiller@starkiller-linux fl_model]$ python drone_client.py
2025-09-18 11:02:41.603617: I tensorflow/core/util/port.cc:153] oneDNN custom operations are on. You may see slightly different numerical results due to floating-point round-off errors from different computation orders. To turn t
, set the environment variable `TF_ENABLE_ONEDNN_OPTS=0`.
2025-09-18 11:02:41.669858: I tensorflow/core/platform/cpu_feature_guard.cc:210] This TensorFlow binary is optimized to use available CPU instructions in performance-critical operations.
To enable the following instructions: SSE3 SSE4.1 SSE4.2 AVX AVX2 AVX_VNNI FMA, in other operations, rebuild TensorFlow with the appropriate compiler flags.
AttributeError: 'MessageFactory' object has no attribute 'GetPrototype'
AttributeError: 'MessageFactory' object has no attribute 'GetPrototype'
AttributeError: 'MessageFactory' object has no attribute 'GetPrototype'
AttributeError: 'MessageFactory' object has no attribute 'GetPrototype'
AttributeError: 'MessageFactory' object has no attribute 'GetPrototype'


--- Starting Local Training on drone_2744 ---
/usr/lib/python3.13/site-packages/keras/src/saving/saving_lib.py:797: UserWarning: Skipping variable loading for optimizer 'adam', because it has 2 variables whereas the saved optimizer has 74 variables.
  saveable.load_own_variables(weights_store.get(inner_path))
Successfully loaded global weights from global_model_initial.weights.h5
Simulating collection of new local data...
Found 30 files belonging to 3 classes.
Starting local fine-tuning...
Epoch 1/5
4/4 ██████████ 13s 129ms/step - accuracy: 0.1667 - loss: 3.5472
Epoch 2/5
4/4 ██████████ 0s 118ms/step - accuracy: 0.1333 - loss: 3.3055
Epoch 3/5
4/4 ██████████ 0s 121ms/step - accuracy: 0.1667 - loss: 2.8302
Epoch 4/5
4/4 ██████████ 0s 116ms/step - accuracy: 0.2000 - loss: 3.0405
Epoch 5/5
4/4 ██████████ 0s 118ms/step - accuracy: 0.0667 - loss: 3.0590
Local fine-tuning complete.
Saved updated weights to drone_2744_update_1758173583.weights.h5

```

```

[starkiller@starkiller-linux fl_model]$ zsh
~/Documents/fl_model > python monthly_aggregation_script.py
2025-09-18 11:01:55.623351: I tensorflow/core/util/port.cc:153] oneDNN custom operations are on. You may see slightly different numerical results due to floating-point round-off errors from different computation orders. To turn t
, set the environment variable `TF_ENABLE_ONEDNN_OPTS=0`.
2025-09-18 11:01:55.689725: I tensorflow/core/platform/cpu_feature_guard.cc:210] This TensorFlow binary is optimized to use available CPU instructions in performance-critical operations.
To enable the following instructions: SSE3 SSE4.1 SSE4.2 AVX AVX2 AVX_VNNI FMA, in other operations, rebuild TensorFlow with the appropriate compiler flags.
AttributeError: 'MessageFactory' object has no attribute 'GetPrototype'
AttributeError: 'MessageFactory' object has no attribute 'GetPrototype'
AttributeError: 'MessageFactory' object has no attribute 'GetPrototype'
AttributeError: 'MessageFactory' object has no attribute 'GetPrototype'
AttributeError: 'MessageFactory' object has no attribute 'GetPrototype'
--- Starting Monthly Federated Averaging ---
Using latest global model as base: global_models/global_model_initial.weights.h5
Found 3 updates to process.
/usr/lib/python3.13/site-packages/keras/src/saving/saving_lib.py:797: UserWarning: Skipping variable loading for optimizer 'rmsprop', because it has 2 variables whereas the saved optimizer has 74 variables.
  saveable.load_own_variables(weights_store.get(inner_path))
Aggregating weights from 3 clients, with a total of 90 new samples.
--- Aggregation Complete ---
New global model saved as: global_models/global_model_2025-09-18.weights.h5
Moved processed updates to drone_updates/processed 2025-09-18

```



**THANK  
YOU!**



[www.reallygreatsite.com](http://www.reallygreatsite.com)

# Results

